

Compilers

2018A7PS0193P

February 8, 2021

1 Process of Translation

The process of compilers converting from source code (the high level language) to target code (the machine language) is known as translation. It consists of many steps, which are described below.

1.1 Lexical Analysis

Lexical analysis follows the following steps:

1. Identify the valid set of characters in the language
2. Break the sequence of characters into appropriate words or tokens (keywords, numbers, operators, etc.)
3. Find out whether these tokens are valid or not.

The key goal of the lexical analyzer is to break a sentence into a series of words/tokens. These breaks are generally done via certain separators. The tokens are recognized via some rules encoded into a Finite State Machine.

During the lexical phase, we may experience the following lexical errors:

- Occurrence of illegal characters
- Exceeding the length of identifier

The output of the lexical analyzer will be a sequence of tokens and their “type”, which signifies whether it is a identifier, operator, etc.

1.2 Syntax Analysis

The syntax analysis takes the sequence of tokens as an input, and generates a parse tree. In case the syntax is not correct according to the grammar rules, it flags a syntactical error. This is modelled using Context Free Grammars that will be recognized using PDAs or Table Driven processes.

1.3 Semantic Analysis

Semantic analysis takes the parse tree as an input, and outputs a disambiguated parse tree. It performs the following:

- Check Semantics
- Error reporting (types, etc.)
- Disambiguate overloaded operators (meaning of operators depends on operands)
- Type coercion (type casting)
- Uniqueness checking (redeclaration of variables)

As such, the disambiguated parse tree gives us an unambiguous representation of the parse tree.

The phases mentioned till now comprise the **front end** of the compiler, where the source code is handled. After this, the compiler works on generating the target code.

1.4 Code optimization

This is an optional phase that modifies the programs to run faster and consume less resources like memory, registers, etc. However, it will not change the representation of the program.

Some examples of machine independent code optimization done is:

- Common sub-expression elimination
- Copy Propagation
- Dead code elimination
- Code Motion
- Strength Reduction

- Constant Folding

TODO: See what these do and examples of each.

1.5 Code Generation

This is the process of mapping from source level abstractions (identifiers, values, etc) to target machine abstractions (registers, memory, etc.). This is a two step process - initially, intermediate code gets generated from the disambiguated parse tree, which is used to generate the final machine code.

During code generation, we have to do the following:

- Map identifiers to locations (memory or registers)
- Map source code operators to opcodes or sequences of opcodes.
- Transform conditionals and iterations to a test/jump or compare instructions
- We use layout parameter passing protocols - the locations for parameters, return values, etc.

1.6 Post Translation Optimizations

Unlike in the code optimization phase where we perform machine independent code optimizations, this does machine-dependent code optimizations. This is an optional phase as well, where we may remove unneeded operations or rearrange to prevent hazards. It is a flexible phase, and may occur at any time in the back-end of the compiler.

1.7 Symbol Table

The symbol table contains information required about the source program identifiers during compilation, including:

- Category of variable
- Data type
- Quantity stored in structure
- Scope information

- Address in Memory

The symbol table must be present in every phase of the compiler, and is used in all the phases to get information about the identifiers.

2 Advantages and Disadvantages of Compilers

The advantages of compilers are:

- Highly modular in nature
- It is retargetable. This means that if there is a single language and multiple machines, then we can use the same front end. If there are many languages
- Source code and machine independent optimizations are possible.

The limitations of the compiler are:

- Design of programming languages has a huge effect on the performance of compilers.
- Lots of work is repeatable. For S languages and M machines, $S \cdot M$ compilers are needed. This is known as the $S * M$ problem of compilers.

The $S * M$ problem is generally solved by introducing some common intermediate language, called the **Universal Intermediate Language Generator**. Some common machine independent intermediate code generation techniques are:

- Postfix Notation
- Three Address code
- Syntax tree
- Directed Acyclic Graph

3 Lexical Analysis

3.1 Functions of the Lexical Analyzer

The lexical analyzer performs the following functions:

- Take high level language as input and output a sequence of tokens
- It generally cleans the code, by stripping off blanks, tabs newlines and comments.
- Keeps track of the line numbers for associated error messages

The lexical analyzer is modelled using regular expressions. As such, it's implementation is done with a DFA. An example of one rule is $L \cdot (L + D)^*$, where L refers to a letter and D refers to a digit.

Definition 3.1. A token is a string of characters which logically belong together, e.g. keywords, number, identifiers, etc.

Definition 3.2. A pattern is the set of strings for which the same token is produced.

Definition 3.3. A lexeme is a sequence of characters matched by a pattern to form the corresponding token.

Now that we understand the definitions, we can see what the lexical analyzer actually does - it transforms strings to the token and passes the lexeme as it's corresponding attribute. For instance, the integer 43 would become `<num,43>`.

3.2 Working of the Lexical Analyzer

The lexical analyzer reads the character one by one from the source code into the lexeme. When it reaches a separator, it assigns a token to the lexeme based on certain rules, and continues to read the characters once more.

However, reading the lexemes character by character is slow, and involves many IO operations. This is done from a buffer instead of directly from the file. Moreover, the prefix of a lexeme is often not enough to determine the token - think of the lexemes `=` and `==`. We instead use a lookahead pointer to determine the appropriate token for a lexeme, and then push back the characters that we do not need in the current lexeme.

3.3 Symbol Table and the Lexical Analyzer

The lexical analyzer also interfaces with the symbol table. When the lexical analyzer discovers a lexeme constituting an identifier, it enters that lexeme to the symbol table. Sometimes, information regarding the token of a particular lexeme may also be store in the symbol table. As such, the symbol table must implement the following operations:

1. `insert(s,t)` : Save lexeme `s` and token `t` and return pointer.

2. `lookup(s)` : return the index of entry for lexeme `s` or '0' if `s` is not found.

To make the symbol table space efficient, we save lexemes in some separate memory, and instead store pointers to the lexemes in the symbol table.

The rule for identifying an identifier and a keyword is generally the same. To be able to tokenize the identifiers and keywords separately, we initialize the symbol table with the list of keywords, say, by calling `insert("if",keyword)`.

3.4 Challenges in Development of Lexical Analyzer

- **Free vs Fixed Lexemes** : A language could specify that lexemes must be in a free or a fixed format. For instance, in a free format, code could look like this.

```
flag = flag
* 6;
```

But in the case of fixed format, this must be entirely in one line. An example of a fixed format language is Python, while a free format language is C.

- **Whitespaces** : How do we deal with whitespaces? Some languages ignore whitespaces until a separator is reached (or interpret contextually), while some languages consider the whitespaces as separators themselves. The former is much more complicated to implement than the latter.
- **Maximal Munch** : The principle of maximal munch directs the lexical analyzer to consume as much available input as possible while creating a construct. This allows us to deal with lexemes like `iff`, and correctly assign it as a identifier rather than the keyword `if`.

3.5 Techniques for specifying tokens

Definition 3.4. Consider R_i is a regular expression and N_i is a unique name, then a regular definition is a series of definitions of the following form

$$N_1 \rightarrow R_1$$

$$N_2 \rightarrow R_2$$

...

$$N_n \rightarrow R_n$$

where each R_i is a regular expression over $\sum \cup \{N_1, N_2, \dots, N_n\}$.

Hence, by assigning a special name N_i to the regular expression R_i , we are in effect defining macros, that remove redundancy in later parts.

The following is an example regular definition for identifiers:

$$\text{Alphabet} \rightarrow A|B|C|\dots|Z|a|b|c|\dots|z$$

$$\text{Digit} \rightarrow 0|1|2|\dots|9$$

$$\text{Identifier} \rightarrow \text{Alphabet}(\text{Alphabet}|\text{Digit})^*$$

This too comes with its own challenges. Regular expressions often fail when identifying the appropriate token, and may pass the invalid tokens to the subsequent translation phases of the compiler (how?). They are only language specifications. Tokenization is a implementation problem.

Tokenization can be done via the following steps:

1. Construct regular expressions for lexemes of each token
2. Construct R matching all lexemes of tokens, so $R = R_1 + R_2 + \dots$, in some well defined precedence order.
3. Consider the input stream to be $S = s_1s_2\dots s_n$. For $i \in [1, n]$, verify whether $s_1\dots s_i \in L(R)$.
4. If $s_1\dots s_i \in L(R) \implies s_1\dots s_i \in L(R_x)$ for some x . We choose the smallest x to be the class of $s_1\dots s_i$.
5. Discard the tokenized input and go back to step 3.

The procedure gives preference to tokens specified earlier using regular expressions. If $s_1\dots s_i \in L(R)$ and $s_1\dots s_j \in L(R)$, we choose the longest prefix, in accordance with the principle of maximal munch.

To implement our regular definitions and recognize tokens, we use **transition diagrams**. They are shown diagrammatically in the same way as Finite Automata. Transitions can be labelled with a symbol, a group of symbols, or regular definitions. A few states may be **retracting states** that indicates that the lexeme does not include the symbol that can bring us to the accepting state.

Sometimes, we may want to push back extra characters into the token stream (think of $>$ and $>=$, we may want to push back the extra character read if the token is $>$). We mark those states with a $*$, to show that we must push back extra characters.

Let us consider the example of hexadecimal and octal constant. The regular definition would be:

$$hex \rightarrow 0|1|2|\dots|9|A|B|C|\dots F$$

$$oct \rightarrow 0|1|2|\dots|7$$

$$Qualifier \rightarrow u|U|l|L$$

$$OctalConstant \rightarrow 0oct^+(Qualifier|\epsilon)$$

$$HexadecimalConstant \rightarrow 0(x|X)hex^+(Qualifier|\epsilon)$$

The transition diagram for the given regular definition is given in Fig 1.

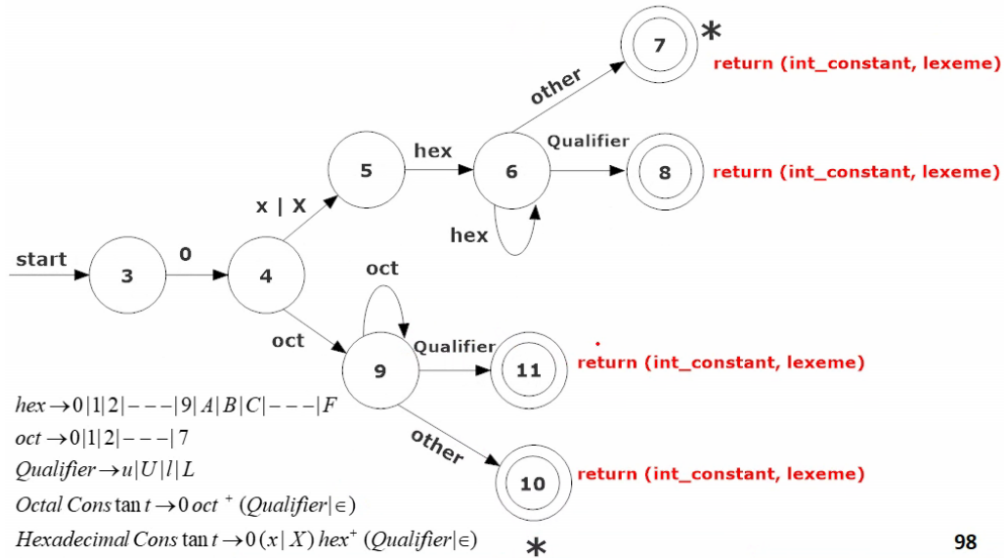


Figure 1: Transition diagram for the Hex and Octal constants

The retracting states are not given. If we get a “bad character”, we should report it as a lexical error.

Let us consider another example of a generalized expression for unsigned numbers. The regular definition is:

$$Digit \rightarrow 0|1|2|\dots|9$$

$$Digits \rightarrow Digit^+$$

$$Fraction \rightarrow ' . ' Digits | \epsilon$$

$$Exponent \rightarrow (E(+|-|\epsilon) Digits) | \epsilon$$

$$Number \rightarrow Digits \cdot Fraction \cdot Exponent$$

The resulting transition diagram from this is in Fig 2.

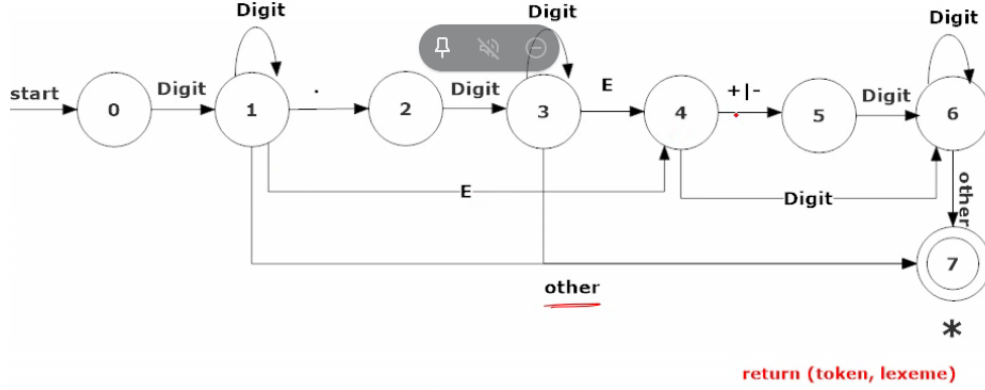


Figure 2: Transition Diagram for Unsigned Numbers

The issue here is this often introduces a lot of complexity. Instead, we may add split this into multiple transition diagrams, to improve implementation complexity. This also empirically better speeds. These multiple transition diagrams can be appropriately combined to generate a lexical analyzer.

The matching process should always start with some transition diagram. If failure occurs in one transition diagram, we retract the forward pointer to the start state, and activate the next transition diagram. If there is failure in all the transition diagrams, we throw a lexical error.

4 Syntax Analysis

4.1 Grammar Rules

Regular definitions are not enough for syntax analysis, since languages are generally not regular. Instead, we use Context Free Grammars to model our language, given by $G = \{T, N, P, S\}$, where

- T is the set of tokens (or terminals)
- N is the set of non-terminals

- P is the set of production rules
- S is the start symbol.

However, just creating the CFG is not the end of our problems - we must choose a method of derivation (left vs right), choose the right non-terminal to expand, choose the right production rule and maintain some precedence order.

When we replace the leftmost non-terminal in any sentential form of the grammar, it is called **leftmost derivation**. Choosing the rightmost non-terminal makes it **rightmost derivation**. If a grammar generates more than one leftmost/rightmost derivation of a string, it is said to be **ambiguous**.

The end goal of our syntax analysis is to create the **parse tree** - where the root is a start symbol, internal nodes are non-terminals and leaf nodes are labelled by tokens. An ambiguous grammar can have more than one parse tree.

Unfortunately, no well defined algorithm exists to remove ambiguity in a grammar. However, we can manually convert any ambiguous grammar into an unambiguous grammar. We can do this by enforcing associativity and precedence rules. This is pretty well covered in PPL and TOC.

In programming languages, ambiguity can come in many forms. For instance, consider the sentence `if A then if B then C1 else C2`. This could be:

```
if A then
{
    if B then
        C1
    else
        C2
}
```

or:

```
if A then
{
    if B then
        C1
    }
else
    C2
```

Since there is ambiguity in its meaning, we must resolve it. One way to do this is to relate an **if** with the closest unrelated **else** (the second interpretation). Another way is to match each **else** with the closest previous **then** (the first interpretation). The unambiguous grammar would then be:

$$\begin{aligned} stmt &\rightarrow matchedstmt | unmatchedstmt \\ matchedstmt &\rightarrow \text{if } exp \text{ then } matchedstmt \text{ else } matchedstmt | others \\ unmatchedstmt &\rightarrow \text{if } exp \text{ then } stmt | \text{if } exp \text{ then } matchedstmt \text{ else } unmatchedstmt \end{aligned}$$

4.2 Parsing

Resolving ambiguity is the problem of the language specification, not of the syntax analysis. Parsing is the job of implementation. We can do this in two ways:

- **Top down parsing:** Construction of the parse tree starting from the root node (start symbol) and proceeds towards the leaves (terminals).
- **Bottom up parsing:** Construction of the parse tree starts from the terminal nodes and builds up towards the root node (start symbol).

4.2.1 Top Down Parsing

In top down parsing, we repeat the following two steps:

1. At a node labelled with non-terminal A , select one of the productions of A and construct the children nodes.
2. Find the next node at which the subtree is to be constructed.

To choose the production rule, as in step 1, backtracking must be done. We try every production rule that we can apply and see which one works. The parser must be intelligent enough to choose the production rule to apply based on the token being pointed to by the pointer.

The **First set**, denoted by $\text{FIRST}(X)$ for a grammar symbol X is the set of tokens that begin the strings derivable from X .

To calculate the First Set, we use the following rules:

- If P is a non terminal and $P \rightarrow Q_1 Q_2 \dots Q_k$ is a production, then if for some i , $\text{First}(Q_i) = \{x\}$ and ϵ is in all of $\text{First}(Q_j)$ (such that $j < i$), then $\text{First}(P) = \{x\}$. If ϵ is in $\text{First}(Q_i)$ for all $i \in [1, k]$, then $\text{First}(P) = \{\epsilon\}$

TODO: Understand and describe First Set better

Let us look at our first parser - the **recursive descent parser**. RDP is a top down method of syntax analysis in which a set of recursive procedures are executed to parse the stream of tokens. A procedure is associated with each non-terminal of the grammar. The procedure generally implements one of the production rules of the grammar.

In each procedure, we perform a match operation on hitting any token on the RHS of the grammar with the current token in the input that needs to be parsed. For example, if we have the string *aba* and the rule $S \rightarrow abB$, it will match the tokens *a* in both the RHS and the string, and move on to the next token. If there is no match, we throw a syntax error.

This approach, of course, has its own limitations. Consider a grammar with two productions:

$$X \rightarrow \gamma_1$$

$$X \rightarrow \gamma_2$$

Suppose $\text{First}(\gamma_1) \cap \text{First}(\gamma_2) \neq \emptyset$, and that *a* is the common terminal symbol. In cases like this, we need to perform backtracking to choose the right production rule. However, RDP does not support backtracking right now. To support it, all productions should be tried in some order. Failure for some productions implies we need to try remaining productions. We would report an error only when there are no other rules.

However, a recursive descent parser may loop forever on productions of the form:

$$A \rightarrow A\alpha$$

This is known as left recursion. We can remove it from the grammar by rewriting the rule as:

$$A \rightarrow \beta A'$$

$$A' \rightarrow \alpha A' | \epsilon$$